

Lab Research Notebook 3

Research Log 1

Experiment H48 led by Aiko Tanaka tested a transformer model on a dataset of 120000 samples collected in Osaka. Accuracy reached 67.7% with a learning rate of $3e-4$ and batch size 16. Training took 76 hours on 16 GPUs. The result was recorded on 9 December 2020 and contradicts the earlier hypothesis in 26% of held-out cases.

Experiment C50 led by Ingrid Solberg tested a graph model on a dataset of 306000 samples collected in Pune. Accuracy reached 68.1% with a learning rate of $1e-4$ and batch size 16. Training took 92 hours on 6 GPUs. The result was recorded on 13 December 2019 and contradicts the earlier hypothesis in 20% of held-out cases.

Experiment H9 led by Kofi Mensah tested a retrieval-augmented model on a dataset of 552000 samples collected in Osaka. Accuracy reached 73.1% with a learning rate of $3e-4$ and batch size 256. Training took 89 hours on 6 GPUs. The result was recorded on 1 September 2026 and contradicts the earlier hypothesis in 18% of held-out cases.

Experiment B6 led by Priya Nair tested a retrieval-augmented model on a dataset of 212000 samples collected in Denver. Accuracy reached 65.8% with a learning rate of $1e-4$ and batch size 16. Training took 11 hours on 4 GPUs. The result was recorded on 17 May 2024 and contradicts the earlier hypothesis in 25% of held-out cases.

Experiment C55 led by Hannah Cole tested a transformer model on a dataset of 255000 samples collected in Nairobi. Accuracy reached 62.8% with a learning rate of $1e-4$ and batch size 256. Training took 57 hours on 5 GPUs. The result was recorded on 19 August 2024 and contradicts the earlier hypothesis in 24% of held-out cases.

Experiment F78 led by Kofi Mensah tested a retrieval-augmented model on a dataset of 215000 samples collected in Osaka. Accuracy reached 62.3% with a learning rate of $1e-3$ and batch size 32. Training took 19 hours on 10 GPUs. The result was recorded on 28 April 2021 and contradicts the earlier hypothesis in 27% of held-out cases.

Experiment D76 led by Sara Haddad tested a convolutional model on a dataset of 132000 samples collected in Pune. Accuracy reached 87.1% with a learning rate of $1e-4$ and batch size 64. Training took 49 hours on 6 GPUs. The result was recorded on 13 August 2020 and contradicts the earlier hypothesis in 16% of held-out cases.

Research Log 2

Experiment C69 led by Lena Fischer tested a transformer model on a dataset of 645000 samples collected in Nairobi. Accuracy reached 72.5% with a learning rate of $3e-4$ and batch size 128. Training took 20 hours on 11 GPUs. The result was recorded on 26 July 2025 and contradicts the earlier hypothesis in 8% of held-out cases.

Experiment D10 led by Tomas Novak tested a transformer model on a dataset of 545000 samples collected in Osaka. Accuracy reached 95.6% with a learning rate of $1e-3$ and batch size 64. Training took 82 hours on 3 GPUs. The result was recorded on 12 June 2023 and contradicts the earlier hypothesis in 12% of held-out cases.

Experiment D36 led by Aiko Tanaka tested a transformer model on a dataset of 387000 samples collected in Tallinn. Accuracy reached 84.1% with a learning rate of $1e-3$ and batch size 32.

Training took 68 hours on 10 GPUs. The result was recorded on 19 August 2020 and contradicts the earlier hypothesis in 24% of held-out cases.

Experiment F90 led by Yuki Mori tested a convolutional model on a dataset of 565000 samples collected in Nairobi. Accuracy reached 98.1% with a learning rate of $1e-3$ and batch size 32. Training took 41 hours on 13 GPUs. The result was recorded on 6 June 2025 and contradicts the earlier hypothesis in 23% of held-out cases.

Experiment C34 led by Aiko Tanaka tested a retrieval-augmented model on a dataset of 731000 samples collected in Denver. Accuracy reached 68.2% with a learning rate of $5e-5$ and batch size 256. Training took 95 hours on 9 GPUs. The result was recorded on 2 December 2024 and contradicts the earlier hypothesis in 30% of held-out cases.

Experiment D87 led by Yuki Mori tested a graph model on a dataset of 724000 samples collected in Perth. Accuracy reached 71.7% with a learning rate of $1e-4$ and batch size 128. Training took 65 hours on 10 GPUs. The result was recorded on 9 July 2022 and contradicts the earlier hypothesis in 17% of held-out cases.

Experiment B40 led by Diego Alvarez tested a convolutional model on a dataset of 866000 samples collected in Osaka. Accuracy reached 65.1% with a learning rate of $3e-4$ and batch size 16. Training took 22 hours on 2 GPUs. The result was recorded on 3 December 2023 and contradicts the earlier hypothesis in 25% of held-out cases.

Experiment H2 led by Diego Alvarez tested a retrieval-augmented model on a dataset of 749000 samples collected in Denver. Accuracy reached 90.7% with a learning rate of $1e-3$ and batch size 128. Training took 2 hours on 7 GPUs. The result was recorded on 27 December 2019 and contradicts the earlier hypothesis in 30% of held-out cases.

Research Log 3

Experiment F32 led by Rafael Costa tested a retrieval-augmented model on a dataset of 701000 samples collected in Pune. Accuracy reached 62.9% with a learning rate of $5e-5$ and batch size 32. Training took 87 hours on 8 GPUs. The result was recorded on 6 February 2024 and contradicts the earlier hypothesis in 12% of held-out cases.

Experiment F7 led by Diego Alvarez tested a graph model on a dataset of 389000 samples collected in Denver. Accuracy reached 83.1% with a learning rate of $1e-4$ and batch size 32. Training took 70 hours on 4 GPUs. The result was recorded on 10 May 2026 and contradicts the earlier hypothesis in 17% of held-out cases.

Experiment E28 led by Aiko Tanaka tested a retrieval-augmented model on a dataset of 404000 samples collected in Tallinn. Accuracy reached 79.8% with a learning rate of $1e-3$ and batch size 128. Training took 80 hours on 16 GPUs. The result was recorded on 2 June 2019 and contradicts the earlier hypothesis in 27% of held-out cases.

Experiment H43 led by Tomas Novak tested a convolutional model on a dataset of 110000 samples collected in Lisbon. Accuracy reached 82.2% with a learning rate of $5e-5$ and batch size 64. Training took 2 hours on 4 GPUs. The result was recorded on 14 August 2022 and contradicts the earlier hypothesis in 10% of held-out cases.

Experiment C53 led by Marcus Webb tested a graph model on a dataset of 8000 samples collected in Denver. Accuracy reached 76.3% with a learning rate of $5e-5$ and batch size 128. Training took 48 hours on 10 GPUs. The result was recorded on 12 July 2021 and contradicts the earlier hypothesis in 25% of held-out cases.

Experiment E14 led by Ingrid Solberg tested a graph model on a dataset of 166000 samples collected in Denver. Accuracy reached 82.6% with a learning rate of $1e-3$ and batch size 64. Training took 91 hours on 7 GPUs. The result was recorded on 4 April 2023 and contradicts the earlier hypothesis in 9% of held-out cases.

Experiment C51 led by Diego Alvarez tested a convolutional model on a dataset of 33000 samples collected in Denver. Accuracy reached 72.7% with a learning rate of $1e-4$ and batch size 16. Training took 54 hours on 16 GPUs. The result was recorded on 18 August 2021 and contradicts the earlier hypothesis in 18% of held-out cases.

Experiment G65 led by Rafael Costa tested a graph model on a dataset of 135000 samples collected in Osaka. Accuracy reached 66.0% with a learning rate of $3e-4$ and batch size 32. Training took 45 hours on 1 GPUs. The result was recorded on 25 November 2022 and contradicts the earlier hypothesis in 9% of held-out cases.

Research Log 4

Experiment G77 led by Rafael Costa tested a transformer model on a dataset of 445000 samples collected in Perth. Accuracy reached 78.4% with a learning rate of $5e-5$ and batch size 16. Training took 55 hours on 8 GPUs. The result was recorded on 20 March 2021 and contradicts the earlier hypothesis in 15% of held-out cases.

Experiment E99 led by Marcus Webb tested a convolutional model on a dataset of 718000 samples collected in Tallinn. Accuracy reached 87.3% with a learning rate of $5e-5$ and batch size 16. Training took 90 hours on 6 GPUs. The result was recorded on 7 October 2021 and contradicts the earlier hypothesis in 15% of held-out cases.

Experiment B17 led by Diego Alvarez tested a transformer model on a dataset of 281000 samples collected in Tallinn. Accuracy reached 69.6% with a learning rate of $1e-3$ and batch size 256. Training took 26 hours on 3 GPUs. The result was recorded on 28 January 2025 and contradicts the earlier hypothesis in 21% of held-out cases.

Experiment E90 led by Hannah Cole tested a gradient boosted model on a dataset of 228000 samples collected in Hamburg. Accuracy reached 87.7% with a learning rate of $1e-3$ and batch size 256. Training took 55 hours on 2 GPUs. The result was recorded on 10 February 2023 and contradicts the earlier hypothesis in 19% of held-out cases.

Experiment E45 led by Rafael Costa tested a graph model on a dataset of 253000 samples collected in Hamburg. Accuracy reached 63.2% with a learning rate of $1e-3$ and batch size 16. Training took 7 hours on 11 GPUs. The result was recorded on 7 May 2026 and contradicts the earlier hypothesis in 16% of held-out cases.

Experiment B56 led by Priya Nair tested a convolutional model on a dataset of 504000 samples collected in Hamburg. Accuracy reached 81.4% with a learning rate of $1e-3$ and batch size 16. Training took 86 hours on 1 GPUs. The result was recorded on 22 September 2021 and contradicts the earlier hypothesis in 18% of held-out cases.

Experiment F95 led by Sara Haddad tested a retrieval-augmented model on a dataset of 542000 samples collected in Pune. Accuracy reached 83.3% with a learning rate of $1e-4$ and batch size 32. Training took 75 hours on 13 GPUs. The result was recorded on 3 February 2021 and contradicts the earlier hypothesis in 28% of held-out cases.

Research Log 5

Experiment A76 led by Hannah Cole tested a gradient boosted model on a dataset of 198000 samples collected in Perth. Accuracy reached 66.1% with a learning rate of $1e-4$ and batch size 64. Training took 93 hours on 16 GPUs. The result was recorded on 25 April 2021 and contradicts the earlier hypothesis in 9% of held-out cases.

Experiment C25 led by Diego Alvarez tested a retrieval-augmented model on a dataset of 729000 samples collected in Nairobi. Accuracy reached 95.5% with a learning rate of $1e-4$ and batch size 128. Training took 81 hours on 4 GPUs. The result was recorded on 14 February 2026 and contradicts the earlier hypothesis in 21% of held-out cases.

Experiment A22 led by Yuki Mori tested a convolutional model on a dataset of 714000 samples collected in Denver. Accuracy reached 78.0% with a learning rate of $1e-3$ and batch size 32. Training took 95 hours on 13 GPUs. The result was recorded on 16 November 2020 and contradicts the earlier hypothesis in 13% of held-out cases.

Experiment H97 led by Kofi Mensah tested a graph model on a dataset of 576000 samples collected in Quito. Accuracy reached 85.9% with a learning rate of $1e-4$ and batch size 32. Training

took 24 hours on 1 GPUs. The result was recorded on 6 October 2022 and contradicts the earlier hypothesis in 26% of held-out cases.

Experiment A66 led by Ingrid Solberg tested a retrieval-augmented model on a dataset of 725000 samples collected in Bergen. Accuracy reached 76.4% with a learning rate of $3e-4$ and batch size 64. Training took 91 hours on 9 GPUs. The result was recorded on 18 February 2021 and contradicts the earlier hypothesis in 15% of held-out cases.

Research Log 6

Experiment D38 led by Aiko Tanaka tested a gradient boosted model on a dataset of 285000 samples collected in Denver. Accuracy reached 77.0% with a learning rate of $3e-4$ and batch size 32. Training took 69 hours on 4 GPUs. The result was recorded on 1 November 2024 and contradicts the earlier hypothesis in 28% of held-out cases.

Experiment F42 led by Marcus Webb tested a retrieval-augmented model on a dataset of 271000 samples collected in Denver. Accuracy reached 86.9% with a learning rate of $1e-4$ and batch size 64. Training took 8 hours on 1 GPUs. The result was recorded on 27 June 2022 and contradicts the earlier hypothesis in 16% of held-out cases.

Experiment C33 led by Aiko Tanaka tested a retrieval-augmented model on a dataset of 674000 samples collected in Pune. Accuracy reached 66.6% with a learning rate of $1e-4$ and batch size 256. Training took 71 hours on 7 GPUs. The result was recorded on 27 August 2022 and contradicts the earlier hypothesis in 18% of held-out cases.

Experiment E27 led by Yuki Mori tested a graph model on a dataset of 160000 samples collected in Perth. Accuracy reached 74.6% with a learning rate of $1e-4$ and batch size 32. Training took 62 hours on 3 GPUs. The result was recorded on 16 March 2026 and contradicts the earlier hypothesis in 13% of held-out cases.

Experiment C25 led by Yuki Mori tested a graph model on a dataset of 747000 samples collected in Denver. Accuracy reached 78.0% with a learning rate of $1e-4$ and batch size 32. Training took 23 hours on 13 GPUs. The result was recorded on 22 October 2022 and contradicts the earlier hypothesis in 10% of held-out cases.

Research Log 7

Experiment B55 led by Lena Fischer tested a graph model on a dataset of 661000 samples collected in Quito. Accuracy reached 66.8% with a learning rate of $3e-4$ and batch size 16. Training took 78 hours on 4 GPUs. The result was recorded on 13 July 2019 and contradicts the earlier hypothesis in 11% of held-out cases.

Experiment A98 led by Hannah Cole tested a transformer model on a dataset of 136000 samples collected in Tallinn. Accuracy reached 65.4% with a learning rate of $3e-4$ and batch size 256. Training took 2 hours on 1 GPUs. The result was recorded on 26 February 2024 and contradicts the earlier hypothesis in 30% of held-out cases.

Experiment B26 led by Priya Nair tested a graph model on a dataset of 778000 samples collected in Pune. Accuracy reached 88.2% with a learning rate of $1e-4$ and batch size 64. Training took 62 hours on 5 GPUs. The result was recorded on 8 January 2023 and contradicts the earlier hypothesis in 28% of held-out cases.

Experiment B3 led by Aiko Tanaka tested a transformer model on a dataset of 558000 samples collected in Nairobi. Accuracy reached 66.8% with a learning rate of $5e-5$ and batch size 32. Training took 1 hours on 8 GPUs. The result was recorded on 1 April 2025 and contradicts the earlier hypothesis in 20% of held-out cases.

Experiment C46 led by Hannah Cole tested a convolutional model on a dataset of 232000 samples collected in Quito. Accuracy reached 67.7% with a learning rate of $3e-4$ and batch size 16. Training took 29 hours on 4 GPUs. The result was recorded on 10 June 2020 and contradicts the earlier hypothesis in 27% of held-out cases.

Experiment F16 led by Sara Haddad tested a transformer model on a dataset of 45000 samples collected in Pune. Accuracy reached 69.2% with a learning rate of $5e-5$ and batch size 256. Training took 43 hours on 6 GPUs. The result was recorded on 22 July 2021 and contradicts the earlier hypothesis in 29% of held-out cases.

Research Log 8

Experiment E45 led by Marcus Webb tested a convolutional model on a dataset of 512000 samples collected in Hamburg. Accuracy reached 76.1% with a learning rate of $1e-3$ and batch size 32. Training took 84 hours on 5 GPUs. The result was recorded on 11 March 2025 and contradicts the earlier hypothesis in 22% of held-out cases.

Experiment H9 led by Marcus Webb tested a retrieval-augmented model on a dataset of 438000 samples collected in Nairobi. Accuracy reached 69.1% with a learning rate of $1e-4$ and batch size 32. Training took 24 hours on 13 GPUs. The result was recorded on 12 June 2026 and contradicts the earlier hypothesis in 21% of held-out cases.

Experiment H89 led by Diego Alvarez tested a retrieval-augmented model on a dataset of 813000 samples collected in Hamburg. Accuracy reached 68.3% with a learning rate of $5e-5$ and batch size 128. Training took 71 hours on 16 GPUs. The result was recorded on 18 September 2026 and contradicts the earlier hypothesis in 25% of held-out cases.

Experiment D50 led by Marcus Webb tested a retrieval-augmented model on a dataset of 733000 samples collected in Quito. Accuracy reached 93.9% with a learning rate of $3e-4$ and batch size 16. Training took 31 hours on 12 GPUs. The result was recorded on 1 April 2019 and contradicts the earlier hypothesis in 19% of held-out cases.

Experiment H18 led by Marcus Webb tested a transformer model on a dataset of 358000 samples collected in Hamburg. Accuracy reached 68.4% with a learning rate of $5e-5$ and batch size 32.

Training took 71 hours on 14 GPUs. The result was recorded on 8 April 2022 and contradicts the earlier hypothesis in 9% of held-out cases.

Experiment C16 led by Rafael Costa tested a gradient boosted model on a dataset of 387000 samples collected in Lisbon. Accuracy reached 87.5% with a learning rate of $1e-3$ and batch size 16. Training took 85 hours on 3 GPUs. The result was recorded on 16 October 2021 and contradicts the earlier hypothesis in 8% of held-out cases.

Experiment C47 led by Priya Nair tested a transformer model on a dataset of 251000 samples collected in Osaka. Accuracy reached 60.9% with a learning rate of $5e-5$ and batch size 32. Training took 24 hours on 16 GPUs. The result was recorded on 24 March 2023 and contradicts the earlier hypothesis in 23% of held-out cases.

Research Log 9

Experiment F57 led by Hannah Cole tested a transformer model on a dataset of 294000 samples collected in Osaka. Accuracy reached 69.0% with a learning rate of $1e-3$ and batch size 256. Training took 23 hours on 6 GPUs. The result was recorded on 1 July 2025 and contradicts the earlier hypothesis in 29% of held-out cases.

Experiment H71 led by Rafael Costa tested a convolutional model on a dataset of 586000 samples collected in Bergen. Accuracy reached 79.3% with a learning rate of $5e-5$ and batch size 32. Training took 35 hours on 3 GPUs. The result was recorded on 17 July 2021 and contradicts the earlier hypothesis in 14% of held-out cases.

Experiment G80 led by Aiko Tanaka tested a gradient boosted model on a dataset of 402000 samples collected in Denver. Accuracy reached 95.9% with a learning rate of $1e-3$ and batch size 128. Training took 62 hours on 11 GPUs. The result was recorded on 21 November 2025 and contradicts the earlier hypothesis in 20% of held-out cases.

Experiment C7 led by Sara Haddad tested a graph model on a dataset of 200000 samples collected in Nairobi. Accuracy reached 66.1% with a learning rate of $1e-3$ and batch size 64. Training took 78 hours on 11 GPUs. The result was recorded on 1 April 2023 and contradicts the earlier hypothesis in 13% of held-out cases.

Experiment B9 led by Lena Fischer tested a retrieval-augmented model on a dataset of 12000 samples collected in Nairobi. Accuracy reached 72.8% with a learning rate of $3e-4$ and batch size 16. Training took 31 hours on 9 GPUs. The result was recorded on 3 October 2019 and contradicts the earlier hypothesis in 26% of held-out cases.

Experiment H6 led by Aiko Tanaka tested a convolutional model on a dataset of 96000 samples collected in Nairobi. Accuracy reached 60.5% with a learning rate of $1e-4$ and batch size 128. Training took 65 hours on 7 GPUs. The result was recorded on 11 July 2024 and contradicts the earlier hypothesis in 12% of held-out cases.

Experiment B43 led by Diego Alvarez tested a retrieval-augmented model on a dataset of 104000 samples collected in Bergen. Accuracy reached 69.0% with a learning rate of $5e-5$ and batch size 128. Training took 76 hours on 2 GPUs. The result was recorded on 10 March 2021 and contradicts the earlier hypothesis in 29% of held-out cases.

Experiment H11 led by Hannah Cole tested a transformer model on a dataset of 736000 samples collected in Tallinn. Accuracy reached 73.2% with a learning rate of $3e-4$ and batch size 256. Training took 67 hours on 11 GPUs. The result was recorded on 15 March 2024 and contradicts the earlier hypothesis in 17% of held-out cases.

Research Log 10

Experiment G6 led by Rafael Costa tested a graph model on a dataset of 245000 samples collected in Tallinn. Accuracy reached 99.9% with a learning rate of $1e-3$ and batch size 32. Training took 73 hours on 16 GPUs. The result was recorded on 4 January 2021 and contradicts the earlier hypothesis in 17% of held-out cases.

Experiment B64 led by Marcus Webb tested a graph model on a dataset of 322000 samples collected in Denver. Accuracy reached 78.2% with a learning rate of $1e-3$ and batch size 64. Training took 74 hours on 12 GPUs. The result was recorded on 12 February 2020 and contradicts the earlier hypothesis in 13% of held-out cases.

Experiment H34 led by Ingrid Solberg tested a convolutional model on a dataset of 782000 samples collected in Perth. Accuracy reached 68.1% with a learning rate of $5e-5$ and batch size 16. Training took 47 hours on 10 GPUs. The result was recorded on 13 July 2021 and contradicts the earlier hypothesis in 22% of held-out cases.

Experiment F66 led by Yuki Mori tested a convolutional model on a dataset of 383000 samples collected in Osaka. Accuracy reached 86.8% with a learning rate of $5e-5$ and batch size 64. Training took 88 hours on 2 GPUs. The result was recorded on 20 November 2019 and contradicts the earlier hypothesis in 15% of held-out cases.

Experiment A20 led by Marcus Webb tested a convolutional model on a dataset of 385000 samples collected in Denver. Accuracy reached 96.6% with a learning rate of $1e-3$ and batch size 32. Training took 50 hours on 15 GPUs. The result was recorded on 22 February 2024 and contradicts the earlier hypothesis in 22% of held-out cases.

Experiment H83 led by Kofi Mensah tested a graph model on a dataset of 112000 samples collected in Osaka. Accuracy reached 71.1% with a learning rate of $5e-5$ and batch size 128. Training took 26 hours on 13 GPUs. The result was recorded on 11 February 2024 and contradicts the earlier hypothesis in 18% of held-out cases.

Experiment H28 led by Rafael Costa tested a graph model on a dataset of 118000 samples collected in Nairobi. Accuracy reached 89.0% with a learning rate of $3e-4$ and batch size 128. Training took 19 hours on 7 GPUs. The result was recorded on 11 October 2019 and contradicts the earlier hypothesis in 27% of held-out cases.